


## Assignment\_12\_Shihan.R

sshihan

Sun Apr 17 10:14:45 2016


```
# # Assignment 12
# # Shaheed Shihan
# # Modern Applied Statistics
#
# 2. We have seen that in  $p = 2$  dimensions, a linear decision boundary
# takes the form  $\theta_0 + \theta_1 X_1 + \theta_2 X_2 = 0$ . We now investigate a non-linear
# decision boundary.
# (a) Sketch the curve
#  $(1 + X_1)^2 + (2 - X_2)^2 = 4$ .

plot(NA, NA, type = "n", xlim = c(-4, 2), ylim = c(-1, 5), asp = 1, xlab = "X
1", ylab = "X2")
symbols(c(-1), c(2), circles = c(2), add = TRUE, inches = FALSE)
```




```
# (b) On your sketch, indicate the set of points for which
#  $(1 + X_1)^2 + (2 - X_2)^2 > 4$ ,
# as well as the set of points for which
#  $(1 + X_1)^2 + (2 - X_2)^2 \leq 4$ .
```

```
plot(NA, NA, type = "n", xlim = c(-4, 2), ylim = c(-1, 5), asp = 1, xlab = "X1", ylab = "X2")
symbols(c(-1), c(2), circles = c(2), add = TRUE, inches = FALSE)
text(c(-1), c(2), "< 4")
text(c(-4), c(2), "> 4")
```



```
# (c) Suppose that a classifier assigns an observation to the blue class
# if
#  $(1 + X_1)^2 + (2 - X_2)^2 > 4$ ,
# and to the red class otherwise. To what class is the observation
#  $(0, 0)$  classified?  $(-1, 1)$ ?  $(2, 2)$ ?  $(3, 8)$ ?
```

```
plot(c(0, -1, 2, 3), c(0, 1, 2, 8), col = c("blue", "red", "blue", "blue"),
     type = "p", asp = 1, xlab = "X1", ylab = "X2")
symbols(c(-1), c(2), circles = c(2), add = TRUE, inches = FALSE)
```



# (d) Argue that while the decision boundary in (c) is not linear in  
 # terms of  $X_1$  and  $X_2$ , it is linear in terms of  $X_1$ ,  $X_2$   
 # 1,  $X_2$ , and  
 #  $X_2^2$

$$\begin{aligned}(1 + X_1)^2 + (2 - X_2)^2 &= 4 \\(X_1)^2 + (X_2)^2 + 2X_1 - 4X_2 + 1 &= 0\end{aligned}$$

which is linear.

# 7. In this problem, you will use support vector approaches in order to  
 # predict whether a given car gets high or low gas mileage based on the  
 # Auto data set.

# (a) Create a binary variable that takes on a 1 for cars with gas  
 # mileage above the median, and a 0 for cars with gas mileage  
 # below the median.

```
library(ISLR)
library(e1071)
```

```
data("Auto")
#head(Auto)
attach(Auto)
```

```

mileage <- as.factor(ifelse(mpg > median(mpg), 1,0))
Auto <- cbind(Auto, mileage)
#head(Auto)
#str(Auto)

# (b) Fit a support vector classifier to the data with various values
# of cost, in order to predict whether a car gets high or low gas
# mileage. Report the cross-validation errors associated with different
# values of this parameter. Comment on your results.

set.seed(1)
tunelinear <- tune(svm,mileage~.,data=Auto,kernel="linear",ranges=list(cost=c
(0.001, 0.01, 0.1, 1,5,10,100)))
summary(tunelinear)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     1
##
## - best performance: 0.01275641
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-03 0.09442308 0.03837365
## 2 1e-02 0.07403846 0.05471525
## 3 1e-01 0.03826923 0.05148114
## 4 1e+00 0.01275641 0.01344780
## 5 5e+00 0.01782051 0.01229997
## 6 1e+01 0.02038462 0.01074682
## 7 1e+02 0.03820513 0.01773427

bestmod1 <- tunelinear$best.model
summary(bestmod1)

##
## Call:
## best.tune(method = svm, train.x = mileage ~ ., data = Auto, ranges = list(
cost = c(0.001,
##   0.01, 0.1, 1, 5, 10, 100)), kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##       cost:  1
##       gamma: 0.003205128

```

```
##
## Number of Support Vectors: 56
##
## ( 26 30 )
##
## Number of Classes: 2
##
## Levels:
## 0 1

#
# (c) Now repeat (b), this time using SVMs with radial and polynomial
# basis kernels, with different values of gamma and degree and
# cost. Comment on your results.

set.seed(1)
tunerad <- tune(svm, mileage~., data=Auto, kernel="radial", ranges=list(cost=c
(0.1,1,10,100,1000),gamma=c(0.5,1,2,3,4)))

summary(tunerad)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
## cost gamma
## 10 0.5
##
## - best performance: 0.04852564
##
## - Detailed performance results:
## cost gamma error dispersion
## 1 1e-01 0.5 0.07916667 0.05201159
## 2 1e+00 0.5 0.05365385 0.05470590
## 3 1e+01 0.5 0.04852564 0.04597747
## 4 1e+02 0.5 0.04852564 0.04597747
## 5 1e+03 0.5 0.04852564 0.04597747
## 6 1e-01 1.0 0.56115385 0.04344202
## 7 1e+00 1.0 0.06634615 0.06187383
## 8 1e+01 1.0 0.06128205 0.06186124
## 9 1e+02 1.0 0.06128205 0.06186124
## 10 1e+03 1.0 0.06128205 0.06186124
## 11 1e-01 2.0 0.56115385 0.04344202
## 12 1e+00 2.0 0.12756410 0.05916990
## 13 1e+01 2.0 0.11730769 0.05277475
## 14 1e+02 2.0 0.11730769 0.05277475
## 15 1e+03 2.0 0.11730769 0.05277475
```

```

## 16 1e-01    3.0 0.56115385 0.04344202
## 17 1e+00    3.0 0.37256410 0.14111555
## 18 1e+01    3.0 0.35205128 0.14165290
## 19 1e+02    3.0 0.35205128 0.14165290
## 20 1e+03    3.0 0.35205128 0.14165290
## 21 1e-01    4.0 0.56115385 0.04344202
## 22 1e+00    4.0 0.48198718 0.04342861
## 23 1e+01    4.0 0.47435897 0.05241275
## 24 1e+02    4.0 0.47435897 0.05241275
## 25 1e+03    4.0 0.47435897 0.05241275

bestmod2 <- tunerad$best.model
summary(bestmod2)

##
## Call:
## best.tune(method = svm, train.x = mileage ~ ., data = Auto, ranges = list(
cost = c(0.1,
##      1, 10, 100, 1000), gamma = c(0.5, 1, 2, 3, 4)), kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##       cost:  10
##       gamma: 0.5
##
## Number of Support Vectors: 259
##
## ( 127 132 )
##
##
## Number of Classes: 2
##
## Levels:
##  0 1

set.seed(1)
tunepol <- tune(svm, mileage ~ ., data = Auto, kernel = "polynomial", ranges
= list(cost = c(0.01, 0.1, 1, 5, 10, 100), degree = c(2, 3, 4)))
summary(tunepol)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost degree
##   100      2
##

```

```

## - best performance: 0.3013462
##
## - Detailed performance results:
##      cost degree      error dispersion
## 1  1e-02      2 0.5611538 0.04344202
## 2  1e-01      2 0.5611538 0.04344202
## 3  1e+00      2 0.5611538 0.04344202
## 4  5e+00      2 0.5611538 0.04344202
## 5  1e+01      2 0.5382051 0.05829238
## 6  1e+02      2 0.3013462 0.09040277
## 7  1e-02      3 0.5611538 0.04344202
## 8  1e-01      3 0.5611538 0.04344202
## 9  1e+00      3 0.5611538 0.04344202
## 10 5e+00      3 0.5611538 0.04344202
## 11 1e+01      3 0.5611538 0.04344202
## 12 1e+02      3 0.3322436 0.11140578
## 13 1e-02      4 0.5611538 0.04344202
## 14 1e-01      4 0.5611538 0.04344202
## 15 1e+00      4 0.5611538 0.04344202
## 16 5e+00      4 0.5611538 0.04344202
## 17 1e+01      4 0.5611538 0.04344202
## 18 1e+02      4 0.5611538 0.04344202

bestmod3 <- tunePol$best.model
summary(bestmod3)

##
## Call:
## best.tune(method = svm, train.x = mileage ~ ., data = Auto, ranges = list(
##   cost = c(0.01,
##     0.1, 1, 5, 10, 100), degree = c(2, 3, 4)), kernel = "polynomial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##     cost:  100
##   degree:  2
##   gamma:  0.003205128
##   coef.0:  0
##
## Number of Support Vectors:  355
##
## ( 177 178 )
##
## Number of Classes:  2
##
## Levels:
##  0 1


```

```
# (d) Make some plots to back up your assertions in (b) and (c).  
# Hint: In the lab, we used the plot() function for svm objects  
# only in cases with p = 2. When p > 2, you can use the plot()  
# function to create plots displaying pairs of variables at a time.  
# Essentially, instead of typing  
# > plot(svmfit , dat)  
# where svmfit contains your fitted model and dat is a data frame  
# containing your data, you can type  
# > plot(svmfit , dat , x1~x4)  
# in order to plot just the first and fourth variables. However, you  
# must replace x1 and x4 with the correct variable names. To find  
# out more, type ?plot.svm.
```


```
fit <- svm(mileage~., data = Auto, kernel = "linear", cost = 1)  
fit2 <- svm(mileage~., data = Auto, kernel = "polynomial", cost = 100, degree  
= 2)  
fit3 <- svm(mileage~. , data = Auto, kernel = "radial", cost = 10, gamma = 0.  
5)  
  
plotp <- function(f){  
  for (name in names(Auto)[!(names(Auto) %in% c("mpg", "mileage", "name"))]) {  
    plot(f, Auto, as.formula(paste("mpg~", name, sep="")))  
  }  
}  
plotp(fit)
```




**SVM classification plot**




**SVM classification plot**




**SVM classification plot**




**SVM classification plot**




**SVM classification plot**



**SVM classification plot**

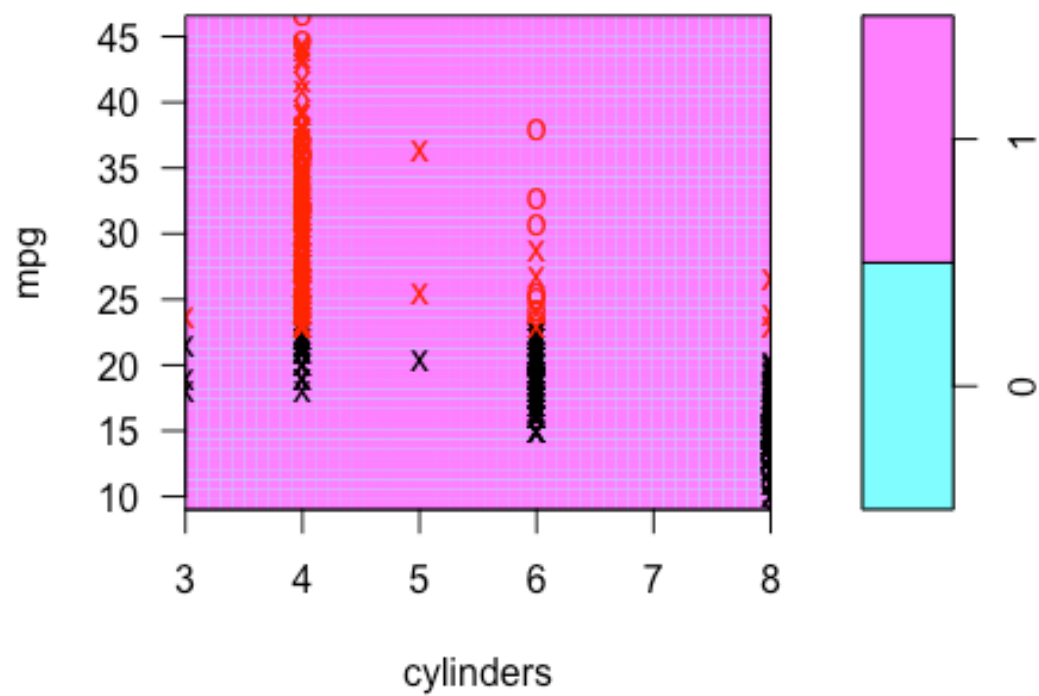


**SVM classification plot**




```
plotp(fit2)
```

**SVM classification plot**







**SVM classification plot**




**SVM classification plot**




**SVM classification plot**




**SVM classification plot**



**SVM classification plot**




**SVM classification plot**




```
plotp(fit3)
```

**SVM classification plot**




**SVM classification plot**







**SVM classification plot**



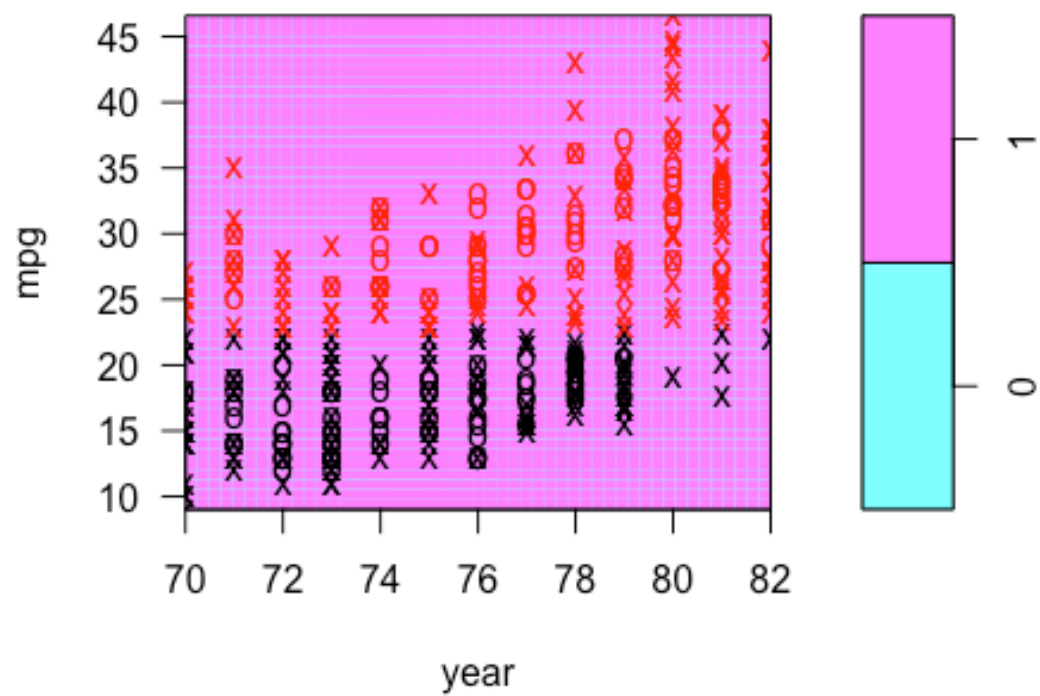
**SVM classification plot**




**SVM classification plot**



**SVM classification plot**



## SVM classification plot



```
detach(Auto)
```

```
# 8. This problem involves the OJ data set which is part of the ISLR  
# package.  
# 372 9. Support Vector Machines  
# (a) Create a training set containing a random sample of 800  
# observations, and a test set containing the remaining  
# observations.
```

```
attach(OJ)
```

```
#head(OJ)
```

```
#dim(OJ)
```

```
train <- sample(nrow(OJ), 800)
```

```
trainOJ <- OJ[train, ]
```

```
testOJ <- OJ[-train, ]
```

```
# (b) Fit a support vector classifier to the training data using  
# cost=0.01, with Purchase as the response and the other variables  
# as predictors. Use the summary() function to produce summary  
# statistics, and describe the results obtained.
```

```
set.seed(4)
```

```

linearsvm <- svm(Purchase~. , data = train0J, kernel = "linear", cost = 0.01)
summary(linearsvm)

##
## Call:
## svm(formula = Purchase ~ ., data = train0J, kernel = "linear",
##      cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:  0.01
##   gamma:  0.05555556
##
## Number of Support Vectors:  448
##
## ( 224 224 )
##
##
## Number of Classes:  2
##
## Levels:
##  CH MM

# At a cost of 0.01, 448 points out of 800 were used
# for creating the support vector. 224 of these belong to MM level and the other 224 to the CH level.

# (c) What are the training and test error rates?
predtrain <- predict(linearsvm, train0J)
t <- table(predtrain, train0J$Purchase)
1-sum(diag(t))/800

## [1] 0.17

##test set
predtest <- predict(linearsvm, test0J)
t1 <- table(predtest, test0J$Purchase)
1-sum(diag(t1))/270

## [1] 0.1518519

# (d) Use the tune() function to select an optimal cost. Consider values
# in the range 0.01 to 10.
set.seed(3)
tunelinear1 <- tune(svm, Purchase~. , data=train0J, kernel="linear", ranges=list(
  cost=c(0.01, 0.05, 0.1, 0.25, 0.5, 0.75,
  1, 5, 10)))
summary(tunelinear1)

```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.1
##
## - best performance: 0.1725
##
## - Detailed performance results:
##   cost   error dispersion
## 1  0.01  0.17625  0.03458584
## 2  0.05  0.17250  0.03717451
## 3  0.10  0.17250  0.03717451
## 4  0.25  0.17500  0.04124790
## 5  0.50  0.17625  0.03928617
## 6  0.75  0.17625  0.04101575
## 7  1.00  0.17750  0.03987829
## 8  5.00  0.17625  0.03557562
## 9 10.00  0.18000  0.03917553

# (e) Compute the training and test error rates using this new value
# for cost.

tunedsvm <- svm(Purchase~. , data = trainOJ, kernel = "linear", cost = 10)
pred1 <- predict(tunedsvm, trainOJ)
t <- table(pred1, trainOJ$Purchase)
1-sum(diag(t))/800

## [1] 0.1625

tunedsvm <- svm(Purchase~. , data = testOJ, kernel = "linear", cost = 10)
pred1 <- predict(tunedsvm, testOJ)
t <- table(pred1, testOJ$Purchase)
1-sum(diag(t))/270

## [1] 0.1259259

# (f) Repeat parts (b) through (e) using a support vector machine
# with a radial kernel. Use the default value for gamma.

set.seed(4)
radialsvm <- svm(Purchase~. , data = trainOJ, kernel = "radial", cost = 0.01)
summary(radialsvm)

##
## Call:
## svm(formula = Purchase ~ ., data = trainOJ, kernel = "radial",
```

```

##      cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##       cost:  0.01
##       gamma: 0.05555556
##
## Number of Support Vectors: 653
##
## ( 325 328 )
##
##
## Number of Classes: 2
##
## Levels:
##  CH MM

# At a cost of 0.01, 653 points out of 800 were used
# for creating the support vector. 328 of these belong to MM level and the other 325 to the CH level.

#training and test error rates?
pred2 <- predict(radialsvm, trainOJ)
t <- table(pred2, trainOJ$Purchase)
1-sum(diag(t))/800

## [1] 0.40625

##test set
predtest <- predict(radialsvm, testOJ)
t1 <- table(predtest, testOJ$Purchase)
1-sum(diag(t1))/270

## [1] 0.3407407

set.seed(3)
tuneradial1 <- tune(svm, Purchase~. , data=trainOJ, kernel="radial", ranges=list(
  cost=c(0.01, 0.05, 0.1, 0.25, 0.5, 0.75,
  1, 5, 10)))
summary(tuneradial1)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     1

```



```

##
## - best performance: 0.17375
##
## - Detailed performance results:
##   cost   error dispersion
## 1  0.01 0.40625 0.05408648
## 2  0.05 0.21250 0.03952847
## 3  0.10 0.18500 0.04556741
## 4  0.25 0.18375 0.04896498
## 5  0.50 0.17875 0.04168749
## 6  0.75 0.17750 0.04158325
## 7  1.00 0.17375 0.04059026
## 8  5.00 0.20000 0.04677072
## 9 10.00 0.19500 0.04571956

tunedsvm <- svm(Purchase~. , data = trainOJ, kernel = "radial", cost = 1)
pred1 <- predict(tunedsvm, trainOJ)
t <- table(pred1, trainOJ$Purchase)
1-sum(diag(t))/800

## [1] 0.16375

pred1 <- predict(tunedsvm, testOJ)
t <- table(pred1, testOJ$Purchase)
1-sum(diag(t))/270

## [1] 0.1518519

# (g) Repeat parts (b) through (e) using a support vector machine
# with a polynomial kernel. Set degree=2.

set.seed(5)
polysvm <- svm(Purchase~. , data = trainOJ, kernel = "polynomial", degree = 2
)
summary(polysvm)

##
## Call:
## svm(formula = Purchase ~ ., data = trainOJ, kernel = "polynomial",
##   degree = 2)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##     cost:    1
##   degree:    2
##   gamma:    0.05555556
##   coef.0:    0
##
## Number of Support Vectors:  463
##

```

```

## ( 229 234 )
##
##
## Number of Classes: 2
##
## Levels:
## CH MM

# At a cost of 1 and degree of 2, 463 points out of 800 were used
# for creating the support vector. 234 of these belong to MM level and the other 229 to the CH level.

#training and test error rates?
pred4 <- predict(polysvm, trainOJ)
t <- table(pred4, trainOJ$Purchase)
1-sum(diag(t))/800

## [1] 0.19

##test set
predtest <- predict(polysvm, testOJ)
t1 <- table(predtest, testOJ$Purchase)
1-sum(diag(t1))/270

## [1] 0.1666667

set.seed(3)
tunepoly1 <- tune(svm, Purchase~. , data=trainOJ, kernel="polynomial", degree =
2, ranges=list(cost=c(0.01, 0.05, 0.1, 0.25, 0.5, 0.75,
1, 5, 10)))
summary(tunepoly1)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
## cost
## 10
##
## - best performance: 0.1825
##
## - Detailed performance results:
## cost error dispersion
## 1 0.01 0.40500 0.05749396
## 2 0.05 0.34625 0.05104804
## 3 0.10 0.33250 0.05407043
## 4 0.25 0.22875 0.03537988
## 5 0.50 0.21500 0.03574602
## 6 0.75 0.20875 0.03283481

```

```
## 7  1.00 0.20625 0.03448530
## 8  5.00 0.18875 0.03793727
## 9 10.00 0.18250 0.03238227

tunedsvm <- svm(Purchase~. , data = trainOJ, kernel = "polynomial", cost = 10
, degree = 2)
pred1 <- predict(tunedsvm, trainOJ)
t <- table(pred1, trainOJ$Purchase)
1-sum(diag(t))/800

## [1] 0.15625

pred1 <- predict(tunedsvm, testOJ)
t <- table(pred1, testOJ$Purchase)
1-sum(diag(t))/270

## [1] 0.1481481

# (h) Overall, which approach seems to give the best results on this data?
# My lowest test error comes from the linear kernel but overall they all seem
# to perform quite well
# with error rates around 16%.
```